COMP 3500 Introduction to Operating Systems
Project 2 – An Introduction to OS/161

Long Version 6.8
Revised: September 11, 2023

Points Possible: 100
Submission via Canvas

There should be no collaboration among students. A student shouldn't share any project code and solution with any other students. Collaborations among students in any form will be treated as a serious violation of the University's academic integrity code.

Programming Environment. You should carry out this project in CentOS 7.9 running on a virtual machine managed by VMware on your local laptop or personal computer. In your CentOS desktop environment, you may make use of Firefox to download the os161 toolchain from Canvas.

Objectives:
- Build `OS/161` and run `Sys/161` on a Linux machine
- Configure and build the `OS/161` kernel
- Discover important design aspects of `OS/161` by examining its source code
- Manage `OS/161` using a version control system called `Git`; apply `Git` and `GitLab` to create a repository and tracking your source code changes
- Use GDB to debug `OS/161`

1. Introduction

The goal of assignment is to provide you with the opportunity to understand the structures of OS/161 and System/161. OS/161 is an operating system on which you will be working, whereas System/161 is a machine simulator on which the OS/161 runs.

This document is organized as follows. The Section 2 outlines Git and GDB – the two programming tools. In Section 3, we will discuss the source code of the OS/161 and the aforementioned tools. Sections "Setting up your account" (see Section 4.1) and its subsequent sections provide instructions on what you have to do for project 2 (also referred to as os161 assignment 0). This project offers you an opportunity to explore the OS/161, which is of crucial importance for projects 3 and 4. Please carefully read and answer the questions in Section 5.2 on page 10. Furthermore, you must submit a small modification to OS/161 (see Section 8 "Practice modifying your kernel" on page 15) and add a few useful debugging hooks to the OS/161.

Although some of the questions allude to concepts we have not yet covered so far, the basic understanding of new concepts can easily be found by skimming the pages in our textbook. Dr. Qin will demonstrate in a lecture how to make use of the `grep` command to consult the code base for answers.

It is desirable to read or skim the entire project specification followed by setting up your account (See Section 4.1 "Setting up your account" on page 4) and working through each section of project 2.

2. OS/161 and System/161

The code you will be working on is comprised of two main components:

- OS/161: a simplified operating system developed by the Systems Research at Harvard group, at Harvard University. You will augment the functionality of the OS/161 in subsequent programming assignments.
- System/161: the machine simulator used to emulate the hardware on which your OS/161 will be running. The focus of this course is to design and implement operating systems rather than developing or simulating hardware. You may not need to change the machine simulator, but are required to modify or augment any portion of the OS/161 code that runs on it.

The OS/161 distribution consists of a full operating system source tree, including some utility programs, libraries, and the like. After you build the OS you boot it on the simulator, analogous to booting a real operating system on real hardware. Using a simulated machine is useful in operating system development, because it is more difficult to develop and debug an operating system running on real and bare hardware. The System/161 machine simulator is an excellent platform for rapid development of operating system code, while retaining a high degree of realism. Apart from floating point support and certain issues relating to RAM cache management, the System/161 provides an accurate emulation of a MIPS R3000 processor.

There will be subsequent OS/161 programming assignments for each of the following topics:
- Synchronization
- Processes and Scheduling
- Memory Management
- File Systems

These projects are cumulative in the sense that you are expected make use of code developed in each assignment to complete subsequent ones. Note that you must reuse your own code to do all assignments, and you have to make sure your code is clean.

3. Tools for Project 2: Git, Gitlab, GDB, and Script

We will learn the following two programming tools that will make your assignments a lot easier. You can find detailed information pertinent to these two tools from the two pdf files: (1) Git-Cheat-Sheet.pdf and (2) gdb-Cheat-Sheet.pdf, which are available on Canvas.

3.1 Git (Version Control System)
Git is a version control system to keep track of changes in source code files. Git is capable of coordinating work on a large number of files developed among multiple software engineers. The main purpose of Git is to manage source code during software development. If you intend to learn more about git, please take a close look at a comprehensive documentation available at: https://www.atlassian.com/git

3.2 GitLab (Web-based Git-Repository Manager)

GitLab is a web-based Git-repository manager with wiki, issue-tracking and CI/CD pipeline features, using an open-source license, developed by GitLab Inc. Access your free GitLab account here: https://gitlab.eng.auburn.edu You can login using your Engineering user name and password.

3.3 GDB (Gnu Debugger)

GDB makes it possible to examine your programs while they are running. Specifically, GDB allows you to execute programs while viewing and setting important values of variables. In particular, GDB enable you to debug OS/161 you are building on the simulated System/161 machine.

Debugging an operating system is similar to debugging ordinary programs. However, a machine has to be reboot when the operating system running on the machine crashes. Fortunately, a machine simulator makes this debugging task easier, because rebooting OS/161 on System/161 takes only a fraction of a second, which is not the case for real operating systems executing on real hardware. More importantly, the System/161 simulator provides information describing what the OS kernel did to cause the crash. Such information may not be easily available when running real OS on real hardware.
To efficiently debug the OS/161, you need to use our version of GDB (referred to as `cs161-gdb`), which is configured for MIPS and has been patched to be able to communicate with your OS kernel through System/161.
A difference between debugging a regular program and debugging an OS/161 kernel is that you need to ensure that you are debugging the operating system rather than the machine simulator. If you were to simply type the obvious thing:

```
cs161-gdb sys161
```

you would be attempting to debug the simulator. For instructions regarding how to debug your OS kernel and for an introduction to GDB, please refer to the gdb.html file (Note: check Canvas for this file).

3.4 The `script` Command

The script command line tool allows you to save a session of your terminal. In addition to saving each command per line in a text file, the script command makes a typescript of everything that happens on your CentOS terminal. Screencasting tools to a desktop session(GUI) is what script is to a terminal. (see also Section 2 in Project 1 Specification).

4. Getting Started

4.1. Setting up Your Account

CAUTION! Please add to your PATH variable the directory containing all of the programs related to OS/161. "-O2" is an O (O for Optimization) not a zero. If you are using bash as your shell you should add the following line near the end of the `~/.bashrc` file.

```
export PATH=~/cs161/bin:$PATH
export CFLAGS="-fcommon -std=gnu89 -g -O2 -Wno-error"
```

If you use `tcsh` as your shell, you should add the following line near the end of the `~/.cshrc`

file:
```
        setenv PATH ~/cs161/bin:$PATH
```
The man page for `tcsh` includes a list of various things you can put in your prompt.

4.2. Getting the Distribution
You have to download, build and install the distributions of the OS/161, System/161 MIPS emulator, and the OS161 toolchain. You will then need to modify some of these instructions to fit the install paths on your machine. The following five compressed files can be found on Canvas->Files->Projects->OS161.

1) tool chain: `cs161-binutils-1.5.tgz`
2) cross compiler: `cs161-gcc-1.5.tgz`
3) special gdb: `cs161-gdb-1.5.tgz`
4) sys161 MIPS emulator: `sys161-1.14.tar.gz`
5) OS/161: `os161-1.10.tar.gz`

CAUTION! Please note that the above sequential order is also the required order for building. Detailed information on building these components is included in the `tgz` files. You may choose to quickly go through a file called "The MIPS toolchain for os161.txt" available on Canvas to learn how to extract and build source code files for the tool chain, cross compiler, and the special gdb. Please read a file named "How to build and run sys161.html" available on Canvas to learn how to compile sys161.

You may follow the instructions below to build the tool chain, cross compiler, special gdb, and sys161. Before starting the following steps, *you need to save the above five tarred and compressed file in `~/cs161`*

Note: If you place the following commands in your ~/.bashrc, you don't have to run the "export" commands in step 4.2.1
```
        export PATH=~/cs161/bin:$PATH
        export CFLAGS="-fcommon -std=gnu89 -g -O2 -Wno-error"
```

4.2.1 Build the tool chain. This process will take approximately two minutes.
```
        export PATH=~/cs161/bin:$PATH
        export CFLAGS="-fcommon -std=gnu89 -g -O2 -Wno-error"
        mkdir ~/cs161
        cd ~/cs161
        tar vfzx cs161-binutils-1.5.tgz
        cd cs161-binutils-1.5
        ./toolbuild.sh
```

4.2.2 Build the cross compiler. It will take about 4 minutes to build the compiler.
```
        cd ~/cs161
        tar vfzx cs161-gcc-1.5.tgz
        cd cs161-gcc-1.5
        ./toolbuild.sh
```

4.2.3 Build the special gdb. It may take you less than two minutes to build the gdb.

```
sudo yum -y install ncurses-devel
cd ~/cs161
tar vfzx cs161-gdb-1.5.tgz
cd cs161-gdb-1.5
./toolbuild.sh
```

If you encounter the "`Configure: error: no termcap library found`" error when you install cs161-gdb-1.5, you may forget to install ncurses-devel. Please read FAQs Question #4 on page 2 for the solution.

4.2.4 Build the sys161 emulator. It will take less than one minute to build the emulator.

```
cd ~/cs161
tar vfzx sys161-1.14.tar.gz
cd sys161-1.14
./configure mipseb
make
make install
```

Note: If you don't complete the above steps (see Sections 4.2.1-4.2.4), you won't be able to compile OS161.

4.3. Scripting Your Session (submit)

4.3.1 Script the following session using the `script` command. This procedure is seemingly a nitpicky item; however, slight deviation from these instructions can produce baffling and time-wasting results - the script may help you to catch subtle mistakes.
```
script
```

4.3.2 Make a directory in which you will do all your programming assignments. Throughout this project, we assume that this directory is called cs161. Moreover, please create a directory called asst0 in which you are going to place the files that you will submit. Note that ~ below is your home directory in Linux.

```
mkdir ~/cs161/asst0
cd ~/cs161
```

4.3.3 Unpack the OS/161 distribution by typing:
```
tar vfzx os161-1.10.tar.gz
```

CAUTION! You must ensure that `os161-1.10.tar.gz` is located in directory `~/cs161` before running the above command to unpack the OS/161 tarball.

4.3.4 The above step creates a directory named `os161-1.10.` Rename your OS/161 source tree to just `os161`.
```
mv os161-1.10 os161
```

4.3.5 End your script session by typing *exit* or by pressing *Ctrl-D*. Rename your typescript file to be `setup.script`.

```
exit

mv typescript ~/cs161/asst0/setup.script
```

4.4   Setting up your GitLab or GitHub repository

CAUTION! You may opt for making use of *GitHub rather than GitLab* to maintain your remote repository.

4.4.1   Use your Engineering account (user name and password) to login to the GitLab website here: https://gitlab.eng.auburn.edu

Note: If you have any problem in accessing your GitLab account, please contact a system administrator at admin@eng.auburn.edu

4.4.2   Create a new empty project; the project name should be os161. The visibility level must be *private*, because you don't want your code to be stolen.

4.4.3   The project created on GitLab will be the remote repository of your project 2. Follow Step 4.5 to create a local repository to be connected to the remote repository constructed in Step 4.4 (i.e., this step).

4.5 Setting up your Git repository (submit)

4.5.1   Script the following session using the script command.
```
script
```

4.5.2   Go to your local workspace.

```
cd ~/cs161/os161
```

4.5.3   Create your Git local repository (i.e., `git init`); connect the local Git repository with your remote one at GitLab (i.e., `git remote add`); commit all the source code files from your workspace (i.e., `~/cs161/os161`) into your local repository (i.e., `git commit`); and push the source code from the local repository to the remote repository at GitLab. You may follow the commands below to complete this step.

```
git init
git add .
git commit -m "Initial OS161 commit"
git branch -M main
git remote add origin
https://gitlab.eng.auburn.edu/usr_name/os161.git
git push -u origin main
```

Note: `usr_name` is your Engineering account's user name. For example, my user name is xzq0001. You must enter "`git remote add origin https://gitlab.eng.auburn.edu/usr_name/os161.git`" as one command without breaking it into two lines.

References:
- Atlassian Git Tutorial

- [Git Branch | Atlassian Git Tutorial](#)

4.5.4   Go to your GitLab project webpage and inspect your remote repository. The URL of your
         GitLab project webpage is listed below, where `usr_name` is your Engineering user
         name (e.g., my user name is xzq0001):
         https://gitlab.eng.auburn.edu/usr_name/os161

4.5.5   Now, you can remove the source tree in your local workspace (i.e., `~/cs161/os161`).

```
cd ~/cs161
rm -rf os161
```

         Don't worry about your deleted local workspace because you have a backup version of
         the source-code tree in your remote GitLab repository. In the next step, you clone a
         copy of the source tree that is yours to work on.

4.5.6   Now, clone a source-code tree from the remote GitLab repository into your local
         machine where the new workspace is the `~/cs161/src dc` directory. `GitLab` will
         ask you to entire your user name and password when you attempt to clone the project.

```
git clone https://gitlab.eng.auburn.edu/usr_name/os161.git src
```

4.5.7   Now use the command below to check your cloned source code tree.

```
ls src
```

4.5.8   End your script session. Rename your script output to `gitinit.script`.

```
exit
mv typescript ~/cs161/asst0/gitinit.script
```

5. Code Reading (submit)

It is challenging that you will be working with a large body of code written by someone else. It is
of importance that you grasp the overall organization of the entire code base, understand
where different pieces of functionality are implemented, and learn how to augment it in a
natural and correct fashion. As you and your partner develop code, although you needn't
understand every detail of your partner's implementation, you still need to understand its
overall structure, how it fits into the greater whole, and how it works.

To become familiar with a code base, you are required to sit down and read the code. You can
use the code reading questions included below to help guide you through reviewing the existing
code. Although there is no need to review every line of code in the system in order to answer all
the questions, you are advised to review every file in the system.
The goal of this exercise is to understand our base system. You should aim at understanding
how it all fits together so that you can make intelligent design decisions when you approach
future assignments. This task is seemingly tedious, but if you understand the structure of the
system, you will have much less difficulty completing future assignments. Although it may not
be apparent yet, you have much more time to study the system's overall structure now than you
will at any other point in the term.

The file system, I/O, and network sections may seem confusing since we have not discussed how these components work. However, it is still useful to review the code now and get a high-level idea of what is happening in each subsystem. It will be fine if you do not understand the low-level details at this point.

The 24 questions below, which appear in red text, are not meant to be tricky--most of the answers can be found in comments in the OS/161 source, though you may have to look elsewhere (such as Silberschatz et al.) for some background information. Place your the answers to the following questions in a file called `~/cs161/asst0/code-reading.txt`.

5.1. Top Level Directory
The top level directory of many software packages is called `src` or `source.` In UNIX, if the operating system source is installed, it is typically found in `/usr/src`. The top of the OS/161 source tree is also called `src.` In this directory, you will find the following files:

`Makefile:` top-level makefile; builds the OS/161 distribution, including all the provided utilities, but does not build the operating system kernel.

`configure:` this is an autoconf-like script. It sets up things like `How to run the compiler.' You needn't understand this file, although we'll ask you to specify certain pathnames and options when you build your own tree.

`defs.mk:` this file is generated when you run `./configure`. You needn't do anything to this file.

`defs.mk.sample:` this is a sample defs.mk file. Ideally, you won't be needing it either, but if configure fails, use the comments in this file to fix `defs.mk.`
and the following directories:

`bin:` this is where the source code lives for all the utilities that are typically found in /bin, e.g., `cat, cp, ls, etc.` The things in bin are considered "fundamental" utilities that the system needs to run.

`include:` these are the include files that you would typically find in `/usr/include` (in our case, a subset of them). These are user level include files; not kernel include files.

`kern:` here is where the kernel source code lives.

`lib:` library code lives here. We have only two libraries: libc, the C standard library, and hostcompat, which is for recompiling OS/161 programs for the host UNIX system. There is also a crt0 directory, which contains the startup code for user programs.

`man:` the OS/161 manual ("man pages") appear here. The man pages document (or specify) every program, every function in the C library, and every system call. The man pages are HTML and can be read with any browser.
`mk:` this directory contains pieces of makefile that are used for building the system. You don't need to worry about these, although in the long run we do recommend that anyone working on

large software systems learn to use make effectively.

`sbin:` this is the source code for the utilities typically found in `/sbin` on a typical UNIX installation. In our case, there are some utilities that let you halt the machine, power it off and reboot it, among other things.

`testbin:` these are pieces of test code.
It is unnecessary to understand every line in every excutable in bin and sbin, but it is worth the time to peruse a couple to see how they work. Eventually, you will want to modify these and/or write your own utilities and these are good models. Similarly, you don't need to read and understand everything in lib and include, but you should know enough about what's there to be able to get around the source tree easily. The rest of this code walk-through is going to concern itself with the `kern` subtree.

5.2 The Kern Subdirectory
Once again, there is a Makefile. This Makefile installs header files but does not build anything. In addition, we have more subdirectories for each component of the kernel as well as some utility directories.

`kern/arch`
This is where architecture-specific code goes. By architecture-specific, we mean the code that differs depending on the hardware platform on which you're running. For our purposes, you need only concern yourself with the `mips` subdirectory.

`kern/arch/mips/conf`
`conf.arch:` This tells the kernel config script where to find the machine-specific, low-level functions it needs (see `kern/arch/mips/mips`).

`Makefile.mips:` Kernel Makefile; this is copied when you "`config a kernel`".
`kern/arch/mips/include`
These files are include files for the machine-specific constants and functions.

1) Which register number is used for the stack pointer (sp) in OS/161?
2) What bus/busses does OS/161 support?
3) What is the difference between `splhigh` and `spl0`?
4) Why do we use typedefs like u_int32_t instead of simply saying "int"?

`kern/arch/mips/mips`
These are the source files containing the machine-dependent code that the kernel needs to run. Most of this code is quite low-level.

5) What does splx return?
6) What is the highest interrupt level?

`kern/asst1`
This is the directory that contains the framework code that you will need to complete assignment 1. You can ignore it for now.

`kern/compile`

This is where you build kernels. In the compile directory, you will find one subdirectory for each kernel you want to build. In a real installation, these will often correspond to things like a debug build, a profiling build, etc. In our world, each build directory will correspond to a programming assignment, e.g., ASST1, ASST2, etc. These directories are created when you configure a kernel (described in the next section). This directory and build organization is typical of UNIX installations but is not universal across all operating systems.

`kern/conf`

config is the script that takes a config file, like ASST0, and creates the corresponding build directory. So, in order to build a kernel, you should:

`cd`

This will create the ASST0 build directory and then actually build a kernel in it. Note that you should specify the complete pathname ./config when you configure OS/161. If you omit the "./", you may end up running the configuration command for the system on which you are building OS/161, and that is almost guaranteed to produce rather strange results!

`kern/dev`

This is where all the low level device management code is stored. For the programming assignments, you can ignore most of this directory.

`kern/include`

These are the include files that the kernel needs. The kern subdirectory contains include files that are visible not only to the operating system itself, but also to user-level programs. (Think about why it's named "kern" and where the files end up when installed.)

7) How frequently are hardclock interrupts generated?
8) What functions comprise the standard interface to a VFS device?
9) How many characters are allowed in a volume name?
10) How many direct blocks does an SFS file have?
11) What is the standard interface to a file system (i.e., what functions must you implement to implement a new file system)?
12) What function puts a thread to sleep?
13) How large are OS/161 pids?
14) What operations can you do on a vnode?
15) What is the maximum path length in OS/161?
16) What is the system call number for a reboot?
17) Where is STDIN_FILENO defined?

`kern/lib`

These are library routines used throughout the kernel, e.g., managing sleep queues, run queues, kernel malloc, etc.

`kern/main`

This is where the kernel is initialized and where the kernel main function is implemented.

`kern/thread`

Threads are the fundamental abstraction on which the kernel is built.

18) Is it OK to initialize the thread system before the scheduler? Why or why not?
19) What is a zombie?
20) How large is the initial run queue?

`kern/userprog`
This is where code is added to create and manage user level processes. As it stands now, OS/161 runs only kernel threads; there is no support for user level code.

`kern/vm`
This directory is also fairly vacant. In Assignment 3, you'll implement virtual memory and most of your code will go in here.

`kern/fs`
The file system implementation has two subdirectories. We'll talk about each in turn.

`kern/fs/vfs`
This is the file-system independent layer (vfs stands for "Virtual File System"). It establishes a framework into which you can add new file systems easily. You will want to go look at vfs.h and vnode.h before looking at this directory.

21) What does a device name in OS/161 look like?
22) What does a raw device name in OS/161 look like?
23) What lock protects the vnode reference count?
24) What device types are currently supported?

`kern/fs/sfs`
This is the simple file system that OS/161 contains by default. You can ignore this directory for now.

6. Building a Kernel (submit)

CAUTION! If you haven't configured your `PATH` variable, you must add the following directory into the `PATH` variable. If you are using bash as your shell you should add the following line near the end of the `~/.bashrc` file.
        `export PATH=~/cs161/bin:$PATH`

Now, we are positioned to build a kernel. As described in See `kern/conf,` you should configure a kernel prior to building it.

6.1 Script the following steps using the `script` command.
        `script`

6.2 Configure your tree for the machine on which you are working. We assume that you work in the directory `~/cs161.` Please note that if you intend to work in a directory that's not `~/cs161` (which you will be doing when you test your later submissions), you will have to

use the `—ostree` option to specify a directory in which you are working. `./configure —help` explains the other options.

```
cd ~/cs161/src
./configure
```

6.3 Configure a kernel named ASST0.

```
cd ~/cs161/src/kern/conf
./config ASST0
```

6.4 Build the ASST0 kernel.

```
cd ../compile/ASST0
make depend
make
```

6.5 Install the ASST0 kernel.

```
make install
```

6.6 Now also build the user level utilties.

```
cd ~/cs161/src
make
```

6.7. End your script session. Rename your script output to `build.script`.

```
exit
mv typescript ~/cs161/asst0/build.script
```


7. Running your kernel (submit)

CAUTION! Download the file `sys161.conf` from Canvas and place it in your OS/161 root directory (`~/cs161/root`).

7.1 Script the following steps using the `script` command.
```
script
```

7.2 Change into your root directory.
```
cd ~/cs161/root
```

7.3 Run the machine simulator on your operating system.
```
./sys161 kernel
```

7.4 At the prompt, type `p /sbin/poweroff <return>`. This tells the kernel to run the "`poweroff`" program that shuts the system down.

7.5 End your script session. Rename your script output to `run.script`.
```
exit
```

```
mv typescript ~/cs161/asst0/run.script
```

8. Practice modifying your kernel (submit)

CAUTION! If you haven't configured your `PATH` variable, you must add the following directory into the `PATH` variable. If you are using bash as your shell you should add the following line near the end of the `~/.bashrc` file.

```
export PATH=~/cs161/bin:$PATH
```

8.1 Script the following session using the `script` command.
```
script
```

8.2 Create a file called `~/cs161/src/kern/main/hello.c`

8.3 Stage the above new file to your local `Git` repository:

```
cd ~/cs161/src/kern/main
git add hello.c
```

8.4 In this file, write a function called `hello()` that uses `kprintf()` to print `"Hello World\n"`.

8.5 Edit `kern/main/main.c` and add a call in a suitable place to `hello()`.

8.6 Edit `kern/conf/conf.kern` using the following commands

```
cd ~/cs161/src/kern/conf
vi conf.kern
/main.c  #search main.c in conf.kern see also line 374
Add the following line:
file  main/hello.c
```

8.7 You must reconfig, and then rebuild (see Section 6 for details)

8.8 Make sure that your new kernel runs and displays the new message. Once your kernel builds, script a session demonstrating the config and build of your modified kernel (see also Sections 6 and 7). Call the output of this script session `newbuild.script.`

```
exit
mv typescript ~/cs161/asst0/newbuild.script
```

9. Using GDB (submit)

You will require two windows for the following portion.

9.1 CAUTION! Script the following GDB session (that is, you needn't script the session in the run window, only the session in the debug window). Be sure both your run window and your debug window are on the same machine.

9.2 Run the kernel in GDB by first running the kernel and then attaching to it from GDB.

```
#In the run window
  cd ~/cs161/root
  ./sys161 -w kernel

#In the debug window
  script
  cd ~/cs161/root
  cs161-gdb kernel
  (gdb) target remote unix:.sockets/gdb
  (gdb) break menu
  (gdb) c
  #gdb will stop at menu()

  (gdb) where

  #displays a nice back trace
  (gdb) detach
  (gdb) quit
```

9.3 End your script session. Rename your script output to gdb.script.
```
  exit
  mv typescript ~/cs161/asst0/gdb.script
```

10. Practice with Git (submit)

CAUTION! You must complete the task specified in Sections 4.4 and 4.5 on page 6 before proceeding to the steps in this Section.

In order to build your kernel above, you already checked out a source tree (see also Sections 4.4 and 4.5). In what follows, we will demonstrate common features of Git. Create a script of the following session (the script should contain everything except the editing sessions; perform those in a different window so they don't appear in the script). Call this file `git-use.script`

10.1 Script the following session using the `script` command.
```
    script
```

10.2 Edit the file `kern/main/main.c.` Add a comment with your name in it.
10.3 Execute
```
    cd ~/cs161/src
    git diff
```
to view the difference between your new version of main.c and the committed one.

References:

- https://www.git-tower.com/learn/git/ebook/en/command-line/advanced-topics/diffs
- https://www.golinuxcloud.com/git-diff-examples/

10.4 Now commit your changes using `git commit.`

10.5 Remove the first 100 lines of `main.c.`

10.6 Try to build your kernel (this ought to fail).

10.7 Realize the error of your ways and get back a good copy of the file.
```
rm main.c
git checkout main.c
```
Reference: https://www.norbauer.com/rails-consulting/notes/git-revert-reset-a-single-file.html

10.8 Try to build your tree again (see Section 6).

10.9 Now, examine the DEBUG macro in `src/kern/include/lib.h`. Based on your earlier reading of the operating system, add ten useful debugging messages to the source code of your operating system os161.
 Note: You shouldn't modify the header file `lib.h`. Rather you must modify any c files (e.g., `main.c`) in the directory of `~/cs161/src`.

A sample source code is given below:
```
DEBUG(DB_VM, "VM free pages: %u\n", free_pages);
```
In the above code, `DB_VM` is a flag defined in `src/kern/include/lib.h`
If the debug DB_VM flag is set, the debug message will be printed on the console.

10.10 Now, display the locations where you inserted these DEBUG statements by doing a diff.
```
cd ~/cs161/src
git diff
```

10.11 Note: Finally, you should create a release using "`git archive`".
```
cd ~/cs161/src
git add .
git commit -m "project 2"
git tag asst0-end
git tag
git show
git push --tags
git archive -o ../asst0/project2_source.tgz HEAD
# Please end your git-use script file here
cd ~/cs161
tar vfzc asst0.tgz asst0
```

Note: The `git archive` command creates a tarball named `project2_source.tgz` in the `~/cs161/asst0` directory. The `tar vfcz` command (i.e., the last command) aims to build a final tarball by compressing `project2_source.tgz` along with the other eight (8) script files in the `~/cs161/asst0` directory (see Step 10.12 below).

Reference: Git Tagging - https://git-scm.com/book/en/v2/Git-Basics-Tagging

Git Archive: https://git-scm.com/docs/git-archive

10.12   End your script session. Rename your script output to `git-use.script`.
```
exit
mv typescript ~/cs161/asst0/git-use.script
```

11. Deliverables
CAUTION! Your `asst0` directory, which you tarred and compressed above, should contain everything you need to submit, specifically:
```
1) setup.script
2) gitinit.script
3) code-reading.txt
4) build.script
5) run.script
6) newbuild.script
7) gdb.script
8) git-use.script
9) project2_source.tgz
```

Now, submit your tarred and compressed file named `asst0.tgz` through Canvas. You must submit your single compressed file through Canvas (no e-mail submission is accepted.

12. Grading Criteria
1) Programing environment setup (i.e., setup.script): 10%
2) Setup Git (i.e., gitinit.script): 10%
3) Reading Code (i.e., code-reading.txt): 30%
4) Build os161 (i.e., build.script): 10%
5) Run os161 (i.e., run.script): 10%
6) Rebuild os161 (i.e., newbuild.script): 10%
7) Use GDB (i.e., gdb.script): 10%
8) Use Git (i.e., cvs-use.script): 10%

13. Late Submission Penalty
- Ten percent (10%) penalty per day for late submission. For example, an assignment submitted after the deadline but up to 1 day (24 hours) late can achieve a maximum of 90% of points allocated for the assignment. An assignment submitted after the deadline but up to 2 days (48 hours) late can achieve a maximum of 80% of points allocated for the assignment.
- Assignment submitted more than 3 days (72 hours) after the deadline will not be graded.

14. Rebuttal Period
- You will be given a period of one week (7 days) to read and respond to the comments and grades of your homework or project assignment. The TA may use this opportunity to address any concern and question you have. The TA also may ask for additional

information from you regarding your homework or project.